

August 30, 2024

```
[ ]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.feature_selection import mutual_info_regression
from sklearn.model_selection import cross_val_score
from xgboost import XGBRegressor

# Set Matplotlib defaults
plt.style.use("seaborn-v0_8-whitegrid")
plt.rc("figure", autolayout=True)
plt.rc(
    "axes",
    labelweight="bold",
    labelsiz="large",
    titleweight="bold",
    titlesiz=14,
    titlepad=10,
)

def plot_variance(pca, width=8, dpi=100):
    # Create figure
    fig, axs = plt.subplots(1, 2)
    n = pca.n_components_
    grid = np.arange(1, n + 1)

    # Explained variance
    evr = pca.explained_variance_ratio_ #
    axs[0].bar(grid, evr) #
    axs[0].set(
        xlabel="Component", title="% Explained Variance", ylim=(0.0, 1.0)
    )

    # Cumulative Variance
    cv = np.cumsum(evr) #
```

```

    axs[1].plot(np.r_[0, grid], np.r_[0, cv], "o-") #
    axs[1].set(
        xlabel="Component", title="% Cumulative Variance", ylim=(0.0, 1.0)
    )

    # Set up figure
    fig.set(figwidth=8, dpi=100)
    return axs

def plot_variance(pca, width=8, dpi=100):
    # Create figure
    fig, axs = plt.subplots(1, 2)
    n = pca.n_components_
    grid = np.arange(1, n + 1)
    # Explained variance
    evr = pca.explained_variance_ratio_
    axs[0].bar(grid, evr)
    axs[0].set(
        xlabel="Component", title="% Explained Variance", ylim=(0.0, 1.0)
    )
    # Cumulative Variance
    cv = np.cumsum(evr)
    axs[1].plot(np.r_[0, grid], np.r_[0, cv], "o-")
    axs[1].set(
        xlabel="Component", title="% Cumulative Variance", ylim=(0.0, 1.0)
    )
    # Set up figure
    fig.set(figwidth=8, dpi=100)
    return axs

def make_mi_scores(X, y):
    X = X.copy()
    for colname in X.select_dtypes(["object", "category"]):
        X[colname], _ = X[colname].factorize()
    # All discrete features should now have integer dtypes
    discrete_features = [pd.api.types.is_integer_dtype(t) for t in X.dtypes]
    mi_scores = mutual_info_regression(X, y,
    ↪discrete_features=discrete_features, random_state=0)
    mi_scores = pd.Series(mi_scores, name="MI Scores", index=X.columns)
    mi_scores = mi_scores.sort_values(ascending=False)
    return mi_scores

def score_dataset(X, y, model=XGBRegressor()):

```

```

# Label encoding for categoricals
for colname in X.select_dtypes(["category", "object"]):
    X[colname], _ = X[colname].factorize()
# Metric for Housing competition is RMSLE (Root Mean Squared Log Error)
score = cross_val_score(
    model, X, y, cv=5, scoring="neg_mean_squared_log_error",
)
score = -1 * score.mean()
score = np.sqrt(score)
return score

def apply_pca(X, standardize=True):
    # Standardize
    if standardize:
        X = (X - X.mean(axis=0)) / X.std(axis=0)
    # Create principal components
    pca = PCA()
    X_pca = pca.fit_transform(X)
    # Convert to dataframe
    component_names = [f"PC{i+1}" for i in range(X_pca.shape[1])]
    X_pca = pd.DataFrame(X_pca, columns=component_names)
    # Create loadings
    loadings = pd.DataFrame(
        pca.components_.T, # transpose the matrix of loadings
        columns=component_names, # so the columns are the principal components
        index=X.columns, # and the rows are the original features
    )
    return pca, X_pca, loadings
df = pd.read_csv("../input/ames.csv")

```

```

[ ]: features = [
    "GarageArea",
    "YearRemodAdd",
    "TotalBsmtSF",
    "GrLivArea",
]

print("Correlation with SalePrice:\n")
print(df[features].corrwith(df.SalePrice))

```

Correlation with SalePrice:

```

GarageArea      0.640138
YearRemodAdd    0.532974
TotalBsmtSF     0.632529
GrLivArea       0.706780
dtype: float64

```

```
[ ]:
```

```
[ ]: X = df.copy()
y = X.pop("SalePrice")
X = X.loc[:, features]

# `apply_pca`, defined above, reproduces the code from the tutorial
pca, X_pca, loadings = apply_pca(X)
print(loadings)
X_pca
```

	PC1	PC2	PC3	PC4
GarageArea	0.541229	-0.102375	-0.038470	0.833733
YearRemodAdd	0.427077	0.886612	-0.049062	-0.170639
TotalBsmtSF	0.510076	-0.360778	-0.666836	-0.406192
GrLivArea	0.514294	-0.270700	0.742592	-0.332837

```
[ ]:
```

	PC1	PC2	PC3	PC4
0	-0.165346	-1.164936	0.233330	0.283527
1	-0.639050	-0.649561	-0.622155	1.740770
2	-0.794227	-1.175790	-0.580254	-0.551055
3	1.636658	-1.907874	-0.675063	-1.052854
4	0.293648	0.610856	0.342374	-0.047752
...	...	...	...	...
2925	-0.276496	0.239255	-0.676658	0.820547
2926	-0.822084	0.414039	-0.593885	0.620322
2927	-1.730457	0.951137	-0.501221	-1.417521
2928	-0.049130	-0.584898	-0.641779	-0.374198
2929	1.090406	0.106617	0.763916	0.328966

[2930 rows x 4 columns]

```
[ ]: X
```

```
[ ]:
```

	GarageArea	YearRemodAdd	TotalBsmtSF	GrLivArea
0	528.0	1960	1080.0	1656.0
1	730.0	1961	882.0	896.0
2	312.0	1958	1329.0	1329.0
3	522.0	1968	2110.0	2110.0
4	482.0	1998	928.0	1629.0
...	...	...	...	...
2925	588.0	1984	1003.0	1003.0
2926	484.0	1983	864.0	902.0
2927	0.0	1992	912.0	970.0
2928	418.0	1975	1389.0	1389.0
2929	650.0	1994	996.0	2000.0

[2930 rows x 4 columns]

```
[ ]: pca.components_
```

```
[ ]: array([[ 0.54122942,  0.42707725,  0.5100756 ,  0.51429428],
          [-0.10237467,  0.88661165, -0.36077798, -0.27069995],
          [-0.03846963, -0.04906151, -0.66683607,  0.74259189],
          [ 0.83373271, -0.17063932, -0.4061918 , -0.33283662]])
```

```
[ ]: X = df.copy()
     y = X.pop("SalePrice")

     X = X.join(X_pca)
     score = score_dataset(X, y)
     print(f"Your score: {score:.5f} RMSLE")
```

Your score: 0.13957 RMSLE

```
[ ]: X
```

```
[ ]: MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  \
0              0         0          141.0  31770.0      0      0          0
1              0         1           80.0  11622.0      0      0          1
2              0         0           81.0  14267.0      0      0          0
3              0         0           93.0  11160.0      0      0          1
4              1         0           74.0  13830.0      0      0          0
...          ...      ...      ...      ...      ...      ...      ...
2925           6         0           37.0   7937.0      0      0          0
2926           0         0            0.0   8885.0      0      0          0
2927           4         0           62.0  10441.0      0      0          1
2928           0         0           77.0  10010.0      0      0          1
2929           1         0           74.0   9627.0      0      0          1

      LandContour  Utilities  LotConfig  ...  MiscFeature  MiscVal  MoSold  \
0                0          0          0  ...          -1         0.0        5
1                0          0          1  ...          -1         0.0        6
2                0          0          0  ...           0  12500.0        6
3                0          0          0  ...          -1         0.0        4
4                0          0          1  ...          -1         0.0        3
...          ...      ...      ...      ...      ...      ...
2925           0          0          2  ...          -1         0.0        3
2926           3          0          1  ...          -1         0.0        6
2927           0          0          1  ...           1       700.0        7
2928           0          0          1  ...          -1         0.0        4
2929           0          0          1  ...          -1         0.0       11

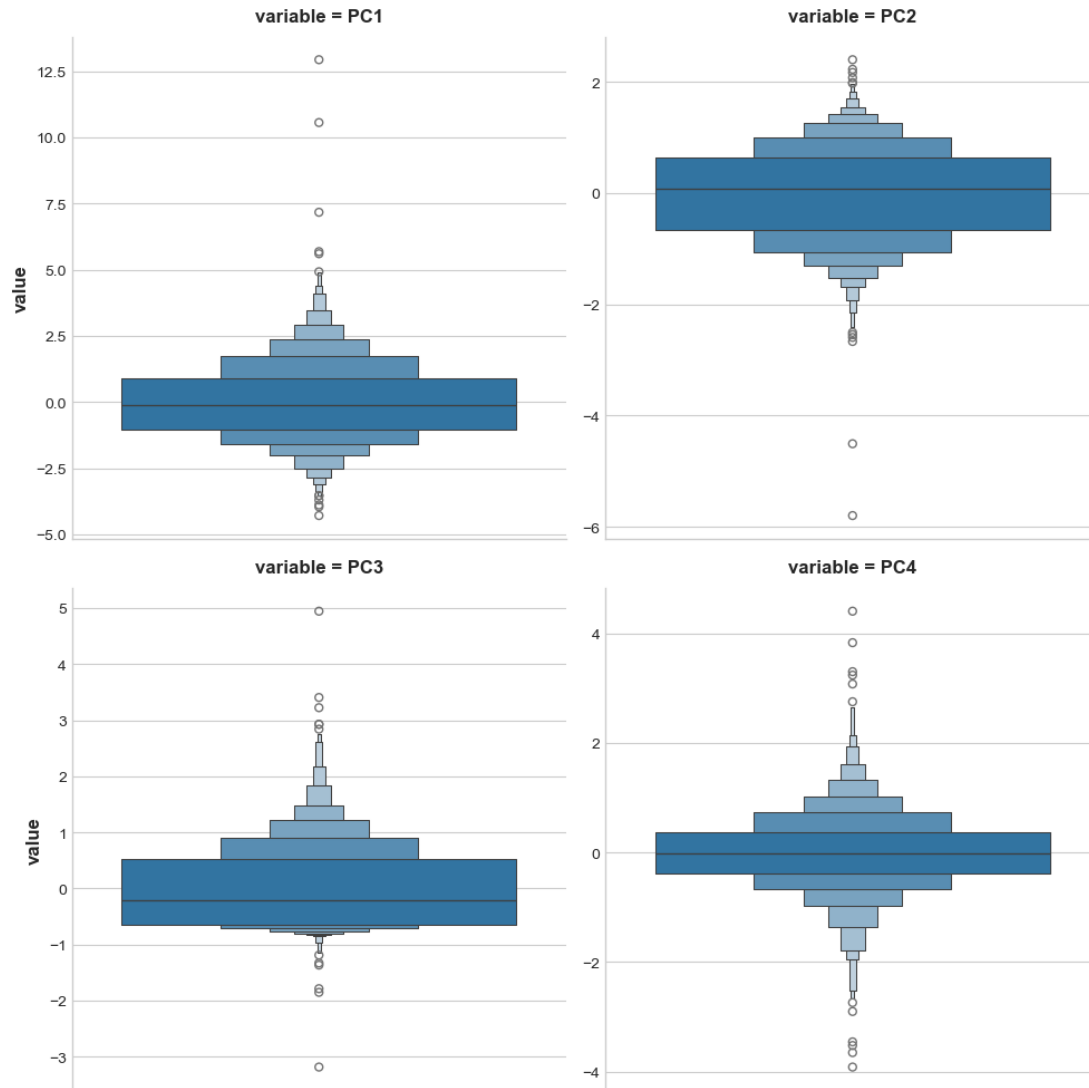
      YearSold  SaleType  SaleCondition      PC1      PC2      PC3  \
0         2010         0              0 -0.165346 -1.164936  0.233330
1         2010         0              0 -0.639050 -0.649561 -0.622155
2         2010         0              0 -0.794227 -1.175790 -0.580254
```

3	2010	0	0	1.636658	-1.907874	-0.675063
4	2010	0	0	0.293648	0.610856	0.342374
...	...	...	...	...	...	...
2925	2006	0	0	-0.276496	0.239255	-0.676658
2926	2006	0	0	-0.822084	0.414039	-0.593885
2927	2006	0	0	-1.730457	0.951137	-0.501221
2928	2006	0	0	-0.049130	-0.584898	-0.641779
2929	2006	0	0	1.090406	0.106617	0.763916

	PC4
0	0.283527
1	1.740770
2	-0.551055
3	-1.052854
4	-0.047752
...	...
2925	0.820547
2926	0.620322
2927	-1.417521
2928	-0.374198
2929	0.328966

[2930 rows x 82 columns]

```
[ ]: sns.catplot(
    y="value",
    col="variable",
    data=X_pca.melt(),
    kind='boxen',
    sharey=False,
    col_wrap=2,
);
```



```
[ ]: # You can change PC1 to PC2, PC3, or PC4
      component = "PC1"

      idx = X_pca[component].sort_values(ascending=False).index
      df.loc[idx, ["SalePrice", "Neighborhood", "SaleCondition"] + features]
```

```
[ ]:      SalePrice      Neighborhood SaleCondition  GarageArea \
      1498      160000      Edwards      Partial      1418.0
      2180      183850      Edwards      Partial      1154.0
      2181      184750      Edwards      Partial      884.0
      1760      745000      Northridge      Abnorml      813.0
      1767      755000      Northridge      Normal      832.0
      ...      ...      ...      ...      ...
```

662	59000	Old_Town	Normal	0.0
2679	80500	Brookside	Normal	0.0
2879	51689	Iowa_DOT_and_Rail_Road	Abnorml	0.0
780	63900	Sawyer	Normal	0.0
1901	39300	Brookside	Normal	0.0

	YearRemodAdd	TotalBsmtSF	GrLivArea
1498	2008	6110.0	5642.0
2180	2009	5095.0	5095.0
2181	2008	3138.0	4676.0
1760	1996	2396.0	4476.0
1767	1995	2444.0	4316.0
...	...	...	...
662	1950	416.0	599.0
2679	1950	0.0	912.0
2879	1950	0.0	729.0
780	1950	0.0	660.0
1901	1950	0.0	334.0

[2930 rows x 7 columns]